

Impacts of the HTTP/2 protocol for large scale web environments


Martin Leucht <martin.leucht@os3.nl>

James Gratchoff <james.gratchoff@os3.nl>

Large Infrastructure Administration
Jaap van Ginkel
Master System and Network Engineering 2015
UvA

Introduction

- HTTP/2 spec on the 18th
- On the way
- 16 years of




I E I F

SG

**Engineering
Steering Group**

Why change?

Total Transfer Size & Total Requests




Source: <http://httparchive.org>

Problems

Bad usage of TCP:
Head of line blocking

Headers size:
Repetition of header



```
Hypertext Transfer Protocol
GET /images/tri-rt-t-14x28.gif HTTP/1.1\r\n
Host: httparchive.org\r\n
Connection: keep-alive\r\n
Pragma: no-cache\r\n
Cache-Control: no-cache\r\n
Accept: image/webp,*/*;q=0.8\r\n
User-Agent: Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/537.36
Referer: http://httparchive.org/\r\n
Accept-Encoding: gzip, deflate, sdch\r\n
Accept-Language: en-US,en;q=0.8,fr;q=0.6\r\n
Cookie: __utmt=1; __utma=108599965.1226105352.1426776704.1426776704
\r\n
```

```
Hypertext Transfer Protocol
GET /images/tri-lft-t-14x28.gif HTTP/1.1\r\n
Host: httparchive.org\r\n
Connection: keep-alive\r\n
Pragma: no-cache\r\n
Cache-Control: no-cache\r\n
Accept: image/webp,*/*;q=0.8\r\n
User-Agent: Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/537.36
Referer: http://httparchive.org/\r\n
Accept-Encoding: gzip, deflate, sdch\r\n
Accept-Language: en-US,en;q=0.8,fr;q=0.6\r\n
Cookie: __utmt=1; __utma=108599965.1226105352.1426776704.1426776704
\r\n
```

Workarounds

- Spriting

TCP works better for larger files. Use one single image for different one.



```
#home {  
  width: 46px;  
  height: 44px;  
  background: url(img_navsprites.gif) 30px 0;  
}  
  
#next {  
  width: 43px;  
  height: 44px;  
  background: url(img_navsprites.gif) 61px 0;  
}
```



- Domain sharding



GET image1.img

GET image2.img

GET image3.img

GET image4.img

Concatenation, Inlining (base 64)...

HTTP/2


Advantages:

- Binary protocol
- one TCP connection
- Streams
- Header compression (HPACK)
- Multiplexing with prioritization and flow control
- Server push
- TLS mandatory (Google Chrome, Firefox)

HTTP/1.1



HTTP/2

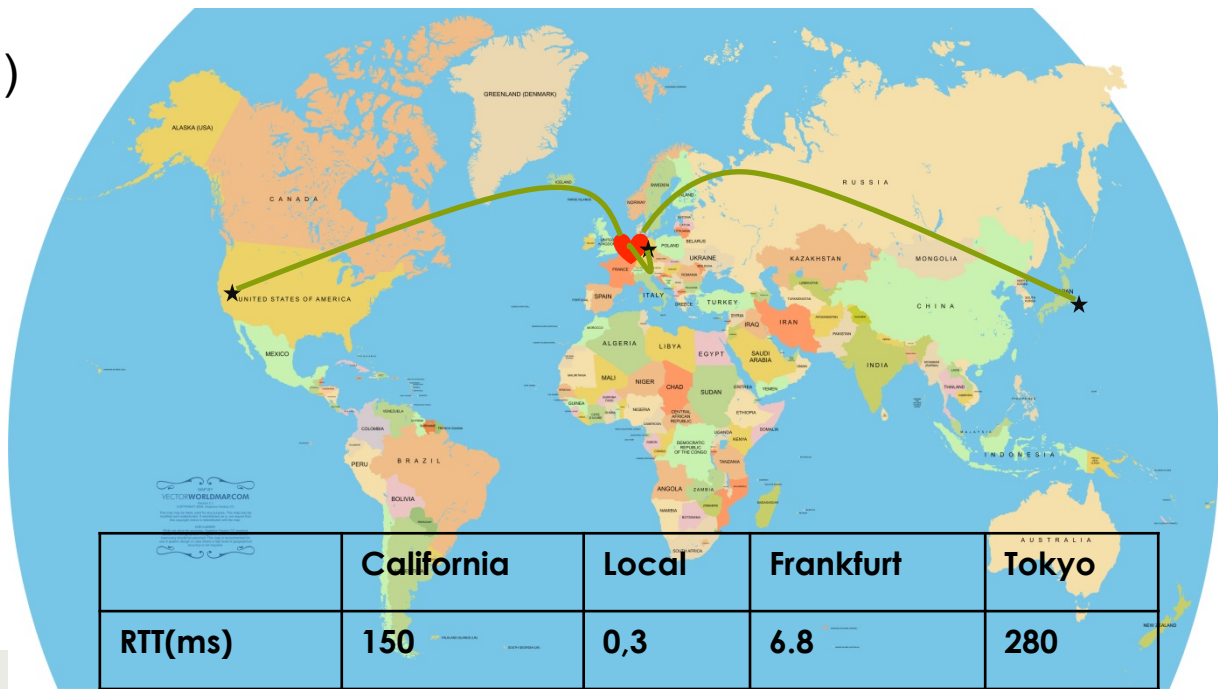


Research questions


- How do the new features of the HTTP/2 protocol improve the performance for frequently visited webpages/webserver?
- What are possible drawbacks that can occur for large web service providers when switching from HTTP/1.1 to HTTP/2? And what is the impact on the infrastructure?
- What is the difference between HTTP/1.1 and HTTP/2 in terms of:
 - Bandwidth and CPU utilization
 - Header size
 - mean time taken for request and response (per location/RTT)

Benchmark setup

- Webserver:
Ubuntu 14 (8CPUx1,8GHz, 8GB RAM)
Servers:
HTTP/2 Nghttp2
HTTP/1.1 Apache2
Reverse Proxy Nghttpx
- 4-Clients:
Ubuntu 14 (ec2-Amazon)
4vCPU, 2GB
Benchmark Tool:
H2load



Httpwatch screenshot



Methods - Measurements

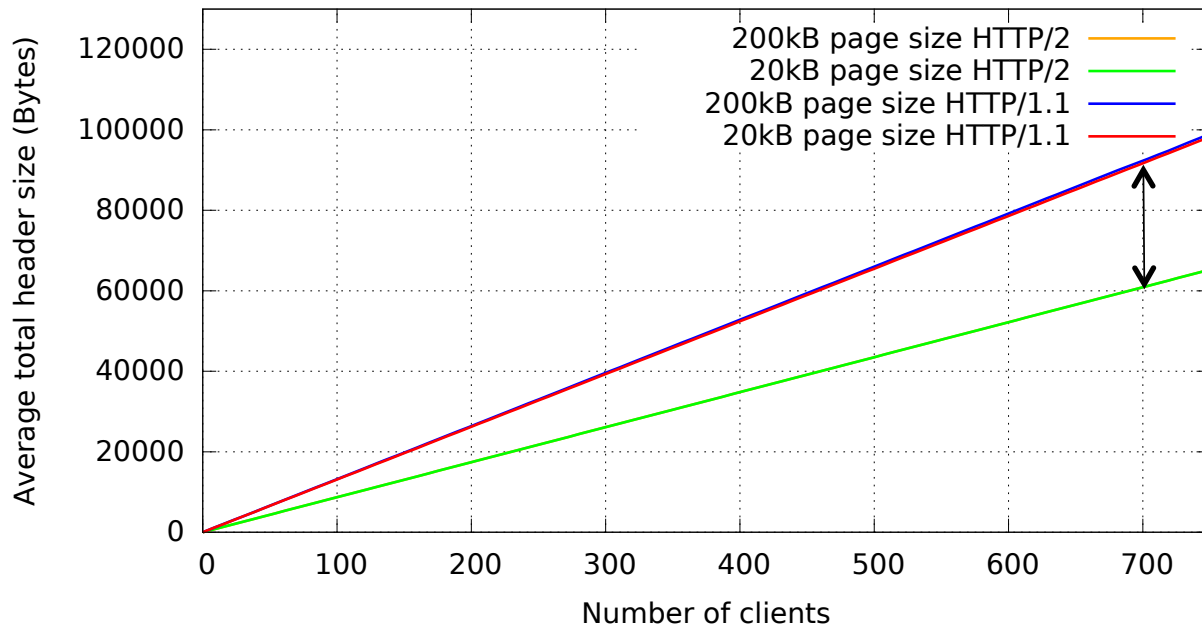
- Test Parameters:
 - HTTP/1.1 and/or HTTP/2 protocol used
 - different locations/distances (RTT)
 - 2 different webpages (size, number of URLs included)
 - increment number of clients/requests in parallel (1...750)
 - recursion depth/repetitions 20

Methods – Data

- Mean of request + response time seen from clients (in ms)
- Header sizes (bytes)
- Failed requests
- Resource Monitoring of web server/reverse proxy:
 - traffic (bps) inbound/outbound
 - CPU utilization
 - number of TCP sockets/connections

Results (HTTP/1.1 vs. HTTP/2 Header Traffic)



Mean header traffic for HTTP/1.1 and HTTP/2 depending on the number of clients and the size of the pages




measured ca. 33% less header traffic in HTTP/2 (header compression)

Results (Frankfurt/Germany – RTT ~7ms)

HTTP/1.1 vs HTTP/2 mean time taken for request and response
depending on the number of
clients and size of pages (RTT ~6.8ms(Frankfurt))




Results (North Carolina/US – RTT ~150ms)



HTTP/2 mean RTT constant (no jitter/packet delay variation)




Results (request time – difference HTTP/2 – HTTP/1.1)



HTTP/2 is in almost all measurements **significantly faster** than HTTP/1.1

Results (HTTP/1.1 and HTTP/2 simultaneous - local)

Mean time taken for request and response for HTTP/1.1 and HTTP/2 clients running in parallel (RTT ~0.2ms(Local))



Conclusion

- Performance improvements HTTP/2 (compared to HTTP/1.1)
 - significantly **less** amount header traffic (ca. 33%)
 - significantly **faster** request/response time – only one TCP session per page request
 - significantly **less** TCP TIME_WAIT connections on server
 - endpoint (server) blocks current connection before it closes it

Conclusion

- Performance improvements HTTP/2 (compared to HTTP/1.1)
 - significantly **less** amount header traffic (ca. 33%)
 - significantly **faster** request/response time – only one TCP session per page request
 - significantly **less** TCP TIME_WAIT connections on server
 - endpoint (server) blocks current connection before it closes it
- Large Scale Webcontent provider should consider when switching to HTTP/2:
 - Use no longer tweaks like 'Spriting' or 'Sharding' for HTTP/2
 - Adapt deep packet inspection firewalls rules (data now binary) in order to allow HTTP/2 requests to path through

Results (Request failures)

Failed Requests comparison between HTTP/1.1 and HTTP/2 requests
from 750 clients and tested from different location

